Supplement: Offline Diagnostics for Direct Elicitation versus Ranked Nodes

ANONYMOUS AUTHOR(S)

1 Scope and procedures

These analyses use the 25 primary targets and saved responses from the completed experiment. They are post hoc and descriptive. No new LLM requests were made. Original predictions, token totals, statistical tests, and reference tables are preserved. The alternative pipelines below are not substituted into the original eight conditions.

The matched direct-anchor variant uses the saved direct response at the exact configurations queried in the paired RNM request, preserving its row order and LLM-selected family. Pairing fixes node, replication, prompt, and evidence condition. Family selection by anchor fit searches WMEAN, WMIN, WMAX, and MIXMINMAX with the original optimizer and chooses the lowest mean anchor TV without reference access. Stronger search retains the original family and uses 16 initial starts plus the incumbent, 80 sweeps per step, and steps 2^{-2} through 2^{-14}. Seed 42, the 10^{-12} improvement threshold, five-point integration, and influence directions are retained. The incumbent remains a candidate. Expanded bounds change mean log ratios to $[-12, 12]$, minimum/maximum weights to $[1, 100]$, and variance to $[10^{-6}, 4]$, retaining the stronger-search candidate.

Reference-informed diagnostics separately fit complete reference tables. Each node considers all four families with stronger search and frozen signs/bounds, then expanded bounds, then all 2^k parent-direction combinations with expanded bounds. Selection minimizes complete-reference TV; these are in-sample best-found approximations, not predictions on hidden configurations or certified global optima. The no-flip expanded candidate is retained in the direction search. Family order breaks exact ties. Numerical plateaus can contain different parameters, so low loss does not establish parameter recovery.

2 Alternative pipelines fitted to LLM judgments

The table uses the original primary non-anchor rows and macro-aggregation: equal targets within networks and equal network versions within the Chrome family; prompt/evidence cells and five replications receive equal weight. The 14 Scrum and 11 Chrome targets remain separate. Uniform probabilities supply a reference-blind uninformed baseline. No additional p-values are calculated.

Pipeline	Scrum TV	Chrome TV
Original direct	0.3184	0.2223
Original RNM	0.2761	0.2517
Uniform	0.4525	0.3739
Direct anchors, paired RNM family	0.2924	0.2456
RNM anchors, family by anchor fit	0.2787	0.2413
RNM anchors, stronger search	0.2761	0.2517
RNM anchors, expanded bounds	0.2761	0.2530

3 Reference-informed approximation diagnostics

The following table reports complete-table TV, including fitted configurations. These values are not comparable to held-out predictive validation or to a mathematical error floor.

Reference-informed search	Scrum TV	Chrome TV
All families, stronger search	0.0560	0.1592
Expanded bounds	0.0559	0.1589
Expanded bounds and direction enumeration	0.0260	0.1322

4 Secondary metrics and replication differences

Metric	Scrum D	Scrum RNM	Chrome D	Chrome RNM
TV	0.3184	0.2761	0.2223	0.2517
Jensen–Shannon (bits)	0.1331	0.1066	0.0905	0.1104
Ordinal CDF	0.1887	0.1588	0.1809	0.2000

The five original replication differences (RNM minus direct TV) are:

Replication	Scrum	Chrome
1	-0.04538	+0.03410
2	-0.03414	+0.02773
3	-0.04732	+0.02682
4	-0.04155	+0.02935
5	-0.04314	+0.02907

5 Illustrative conditional distributions

For each stratum, we select the target nearest the unweighted median per-target RNM-minus-direct non-anchor TV effect, breaking ties by identifier. We show its all-middle aligned anchor and average the 20 saved distributions per method. This transparent rule does not select the largest benefit or harm. Mean probability vectors are illustrative; their TV need not equal mean trial TV.

5.1 All contributors

Number of distinct component contributors.

Parents: Component changes = [11, 51); Distinct reviewers = [3, 7).

Child state order: [1, 3), [3, 5), [5,102]. The target's mean non-anchor effect is +0.0193. All 20 RNM trials select WMEAN.

Distribution	State 1	State 2	State 3
Reference	0.356	0.421	0.223
Direct	0.261	0.358	0.382
Raw RNM	0.213	0.396	0.391
Fitted RNM	0.231	0.537	0.231

At this row, mean trial TV is direct: 0.1587, raw rnm: 0.1691, fitted rnm: 0.1318. These anchor illustrations do not establish a universal pattern for non-anchor rows.

5.2 Team Orientation

Mutual trust, respect and motivation to prioritize team goals over isolated individual goals.

Parents: Personal Attributes = Medium; Expertise = Medium.

Child state order: Low, Medium, High. The target's mean non-anchor effect is −0.0306. All 20 RNM trials select WMEAN.

Distribution	State 1	State 2	State 3
Reference	0.053	0.927	0.020
Direct	0.237	0.501	0.263
Raw RNM	0.206	0.585	0.210
Fitted RNM	0.207	0.586	0.207

At this row, mean trial TV is direct: 0.4260, raw rnm: 0.3425, fitted rnm: 0.3414. These anchor illustrations do not establish a universal pattern for non-anchor rows.

6 Qualitative evidence packets

All 61 retained nodes have nonempty packets containing 1–3 cards. Card text totals 24–103 whitespace-separated words per node, with median 72. The same node packet is reused across methods, prompt variants, and replications. The following are the complete cards for the two illustrative targets. Both evidence arms also receive the child and parent definitions.

6.1 All contributors

Component history. Changes are aggregated at component-directory level. Size, change volume and contributor counts describe different aspects of development activity; one does not uniquely determine another.

Contributor roles. Contributors are classified by their share of component changes. Minor and major contributors partition the contributor set. Holding total contributors fixed, a larger minor-contributor count leaves fewer major contributors.

Review measures. Review activity can include discussion, approval and revisions of a patch. A patch revision during review is not itself a confirmed post-release defect. Author-only approval and review without discussion describe distinct practices.

6.2 Team Orientation

Mutual support. Team orientation involves mutual respect and giving team goals priority. Constructive relationships and confidence in colleagues help members work together toward those goals.

Knowledge coverage. Expertise includes collective knowledge and overlapping capabilities. Being able to perform parts of another member's work supports flexibility when circumstances demand coverage.

Assessment scope. The assessment is intended to support retrospective discussion. Project-specific practices and metrics can be added to a general model. A contextual ordinal rating does not have a universal numerical threshold across organizations.

7 Target-level primary differences

7.1 Scrum

Network	Target	Direct	RNM	Δ
Scrum	Teamwork Quality	0.3050	0.3095	+0.0044
Scrum	Increment quality	0.3306	0.3510	+0.0204
Scrum	Customer satisfaction	0.1131	0.0694	-0.0437
Scrum	Cohesion	0.2471	0.1018	-0.1452
Scrum	Self-Organizing	0.3491	0.3487	-0.0004
Scrum	Collaboration	0.3671	0.2777	-0.0894
Scrum	Team Orientation	0.3869	0.3563	-0.0306
Scrum	Communication	0.4366	0.3707	-0.0660
Scrum	Coordination	0.1500	0.1284	-0.0216
Scrum	Continuous backlog refinement	0.2454	0.1708	-0.0746
Scrum	Sprint execution quality	0.3522	0.3316	-0.0206
Scrum	Requirements volatility	0.5307	0.4614	-0.0692
Scrum	Product backlog quality	0.3366	0.3062	-0.0304
Scrum	Sprint planning quality	0.3074	0.2821	-0.0253

7.2 Chrome

Network	Target	Direct	RNM	Δ
chrome_40	Distinct reviewers	0.1189	0.1419	+0.0229
chrome_40	Major contributors	0.0676	0.0728	+0.0052
chrome_40	Minor contributors	0.3389	0.3860	+0.0471
chrome_40	Prior defects	0.1728	0.1240	-0.0487
chrome_40	Patch revisions during review	0.3301	0.3146	-0.0156
chrome_40	Self-approved changes	0.2133	0.2219	+0.0086
chrome_42	All contributors	0.2139	0.2332	+0.0193
chrome_42	Major-author ownership	0.4344	0.4255	-0.0089
chrome_42	Major contributors	0.0000	0.1565	+0.1565
chrome_42	Minor contributors	0.3696	0.4561	+0.0866
chrome_42	Self-approved changes	0.1706	0.1949	+0.0243

These target means describe this sample; targets and rows are not independent experimental replications.